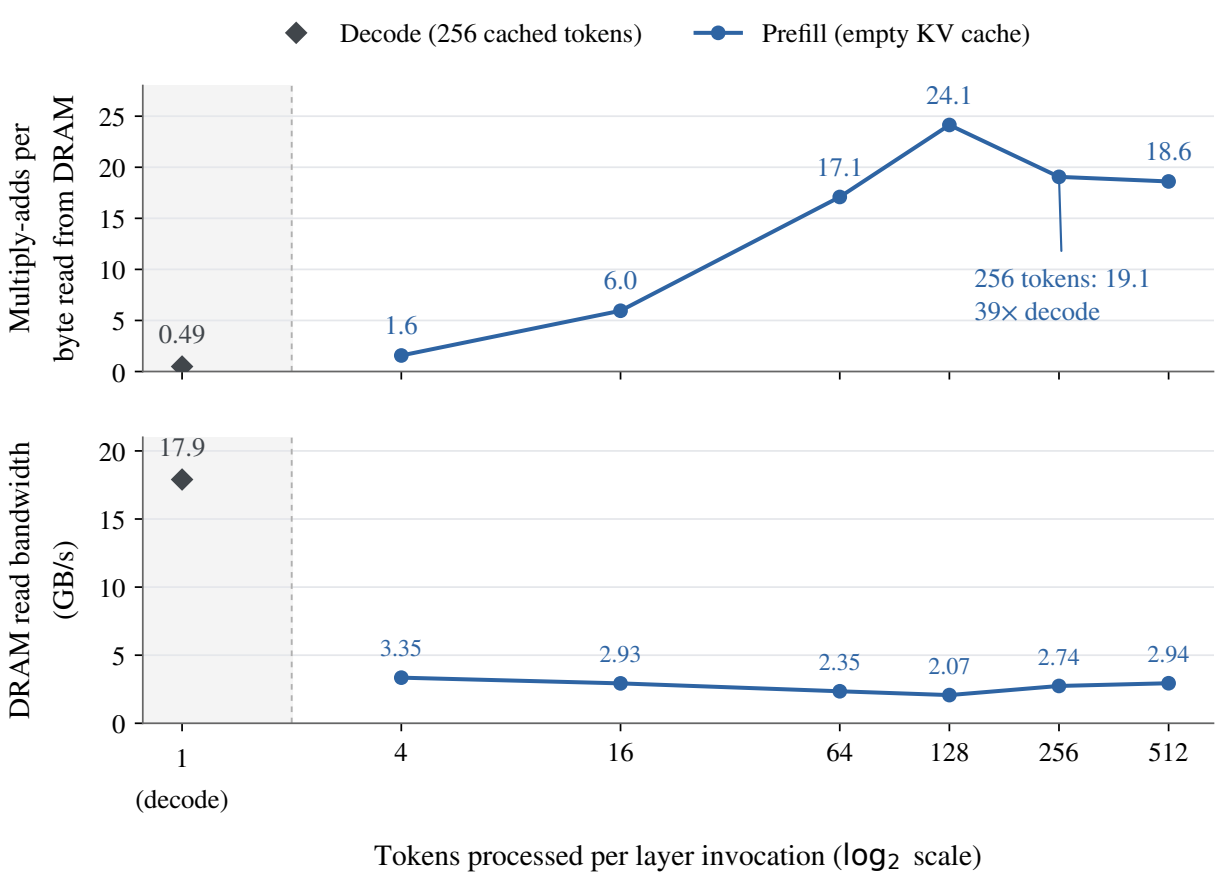
Decode (256 cached tokens)
Prefill (empty KV cache)
Multiply-adds per byte read from DRAM
0
5
10
15
20
25
0.49
1.6
6.0
17.1
24.1
18.6
256 tokens: 19.1
39× decode
DRAM read bandwidth (GB/s)
0
5
10
15
20
17.9
3.35
2.93
2.35
2.07
2.74
2.94
1
(decode)
4
16
64
128
256
512
Tokens processed per layer invocation (log2 scale)