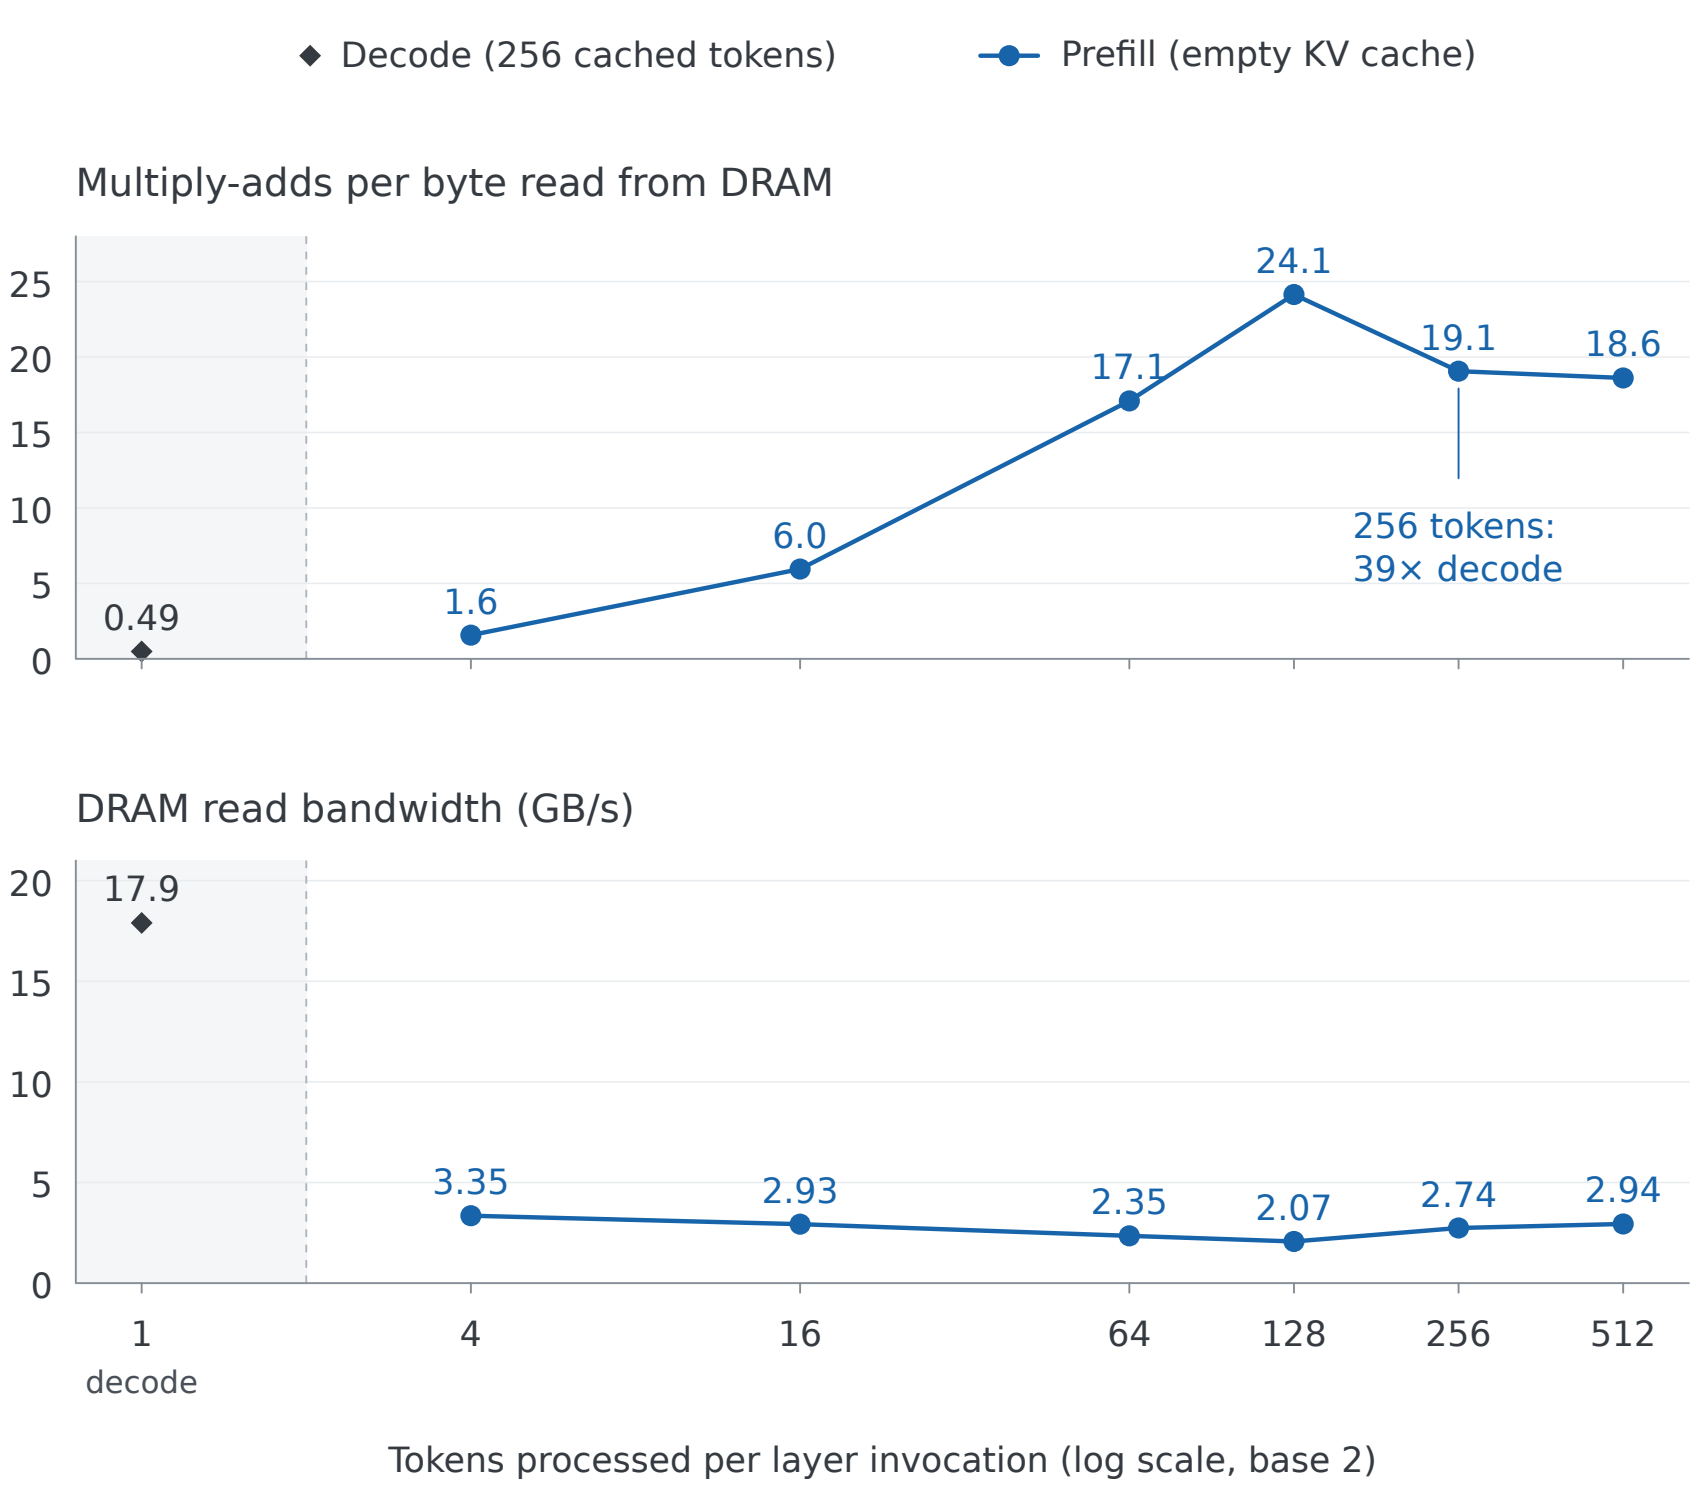
Decode (256 cached tokens)
Prefill (empty KV cache)
Multiply-adds per byte read from DRAM
25
20
15
10
5
0
0.49
1.6
6.0
17.1
24.1
19.1
18.6
256 tokens:
39× decode
DRAM read bandwidth (GB/s)
20
15
10
5
0
17.9
3.35
2.93
2.35
2.07
2.74
2.94
1
decode
4
16
64
128
256
512
Tokens processed per layer invocation (log scale, base 2)